

```

1 # Cracking the Market Code with AI-Driven Stock Price Prediction Using Time Series Analysis
2
3 # =====
4 # 1. Importing Required Libraries
5 # =====
6 import numpy as np
7 import pandas as pd
8 import matplotlib.pyplot as plt
9 import seaborn as sns
10 from sklearn.preprocessing import MinMaxScaler
11 from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
12 from tensorflow.keras.models import Sequential
13 from tensorflow.keras.layers import LSTM, Dense, Dropout
14 import yfinance as yf
15
16 # =====
17 # 2. Data Collection
18 # =====
19 data = yf.download('AAPL', start='2020-01-01', end='2023-12-31')
20 data.to_csv("apple_stock.csv")
21
22 # =====
23 # 3. Data Preprocessing
24 # =====
25 data.dropna(inplace=True)
26 data = data[['Open', 'High', 'Low', 'Close', 'Volume']]
27
28 scaler = MinMaxScaler()
29 data_scaled = scaler.fit_transform(data)
30
31 # =====
32 # 4. Feature Engineering
33 # =====
34 X = []
35 y = []
36 window = 60
37 for i in range(window, len(data_scaled)):
38     X.append(data_scaled[i-window:i])
39     y.append(data_scaled[i, 3]) # Predicting 'Close'
40
41 X, y = np.array(X), np.array(y)
42
43 # =====
44 # 5. Train-Test Split
45 # =====
46 split = int(0.8 * len(X))
47 X_train, X_test = X[:split], X[split:]
48 y_train, y_test = y[:split], y[split:]
49
50 # =====
51 # 6. Model Building (LSTM)
52 # =====
53 model = Sequential()
54 model.add(LSTM(units=50, return_sequences=True, input_shape=(X.shape[1], X.shape[2])))
55 model.add(Dropout(0.2))
56 model.add(LSTM(units=50))
57 model.add(Dropout(0.2))
58 model.add(Dense(1))
59 model.compile(optimizer='adam', loss='mean_squared_error')
60
61 # =====
62 # 7. Model Training
63 # =====
64 model.fit(X_train, y_train, epochs=20, batch_size=32, validation_data=(X_test, y_test))
65
66 # =====
67 # 8. Evaluation
68 # =====
69 y_pred = model.predict(X_test)
70 y_test_inv = scaler.inverse_transform(np.concatenate((np.zeros((y_test.shape[0], 3)),
71                                                         np.expand_dims(y_test, axis=1),
72                                                         np.zeros((y_test.shape[0], 1))), axis=1))[:, 3]
73 y_pred_inv = scaler.inverse_transform(np.concatenate((np.zeros((y_pred.shape[0], 3)),
74                                                         y_pred,
75                                                         np.zeros((y_pred.shape[0], 1))), axis=1))[:, 3]
76
77 print("MAE:", mean_absolute_error(y_test_inv, y_pred_inv))
78 print("RMSE:", np.sqrt(mean_squared_error(y_test_inv, y_pred_inv)))
79 print("R2 Score:", r2_score(y_test_inv, y_pred_inv))
80
81 # =====
82 # 9. Visualization
83 # =====
84 plt.figure(figsize=(12, 6))
85 plt.plot(y_test_inv, label="Actual")
86 plt.plot(y_pred_inv, label="Predicted")
87 plt.title("Actual vs Predicted Stock Price")
88 plt.xlabel("Time")
89 plt.ylabel("Price")
90 plt.legend()
91 plt.tight_layout()
92 plt.show()
93
94 # =====
95 # 10. Save Model
96 # =====
97 model.save("lstm_stock_model.h5")
98

```